EXPLORATORY QUALITATIVE EVIDENCE MAP

The Plural of Anecdote Is Data

How Patients Describe Using Large Language Models for Health

Birger Moëll

Abstract

Large language models (LLMs) are increasingly used by patients and caregivers for health-related sensemaking outside formal clinical settings. Existing research has largely evaluated LLM performance on medical benchmarks, clinician-authored prompts, or simulated patient questions. Less is known about how people describe using these systems in everyday illness, uncertainty, and care navigation. We compiled a seed corpus of 100 public online accounts in which patients, caregivers, family members, or peers described using LLMs or chatbot-like AI systems for health-related purposes. Sources included Reddit posts and comments, patient forums, personal essays, news-mediated patient accounts, disease-specific publications, and case-report-adjacent sources. We pilot-coded 35 strategically selected records by source quality, account type, health domain, primary use case, reported outcome, evidence quality, system gap, and risk mechanism. Across the corpus, LLMs were most often described as useful when the task was sensemaking: translating medical language, organizing longitudinal records, preparing questions for clinicians, identifying possible patterns, supporting behavior change, or providing emotional containment between encounters. These accounts suggest that patients are using LLMs less as stand-alone doctors than as always-available tools for the invisible work around healthcare. However, the same properties that make LLMs useful -- fluency, patience, personalization, and availability -- also appear in harm narratives involving wrong triage, unsafe medication or diet advice, compulsive reassurance seeking, psychosis or mania amplification, model-dependent coping, and privacy exposure. Public anecdotes cannot establish clinical efficacy, but they are data about patient behavior, perceived unmet need, and failure modes that may be missed by benchmark studies. We argue that anecdotal corpora can help researchers, clinicians, and system designers identify where LLMs amplify care, where they compensate for system gaps, and where they quietly place patients at risk.

Keywords

large language models; ChatGPT; patient narratives; digital health; qualitative research; health information seeking; anecdotal evidence; patient safety

Introduction

Large language models have entered health behavior faster than clinical governance can evaluate them. Patients now ask chatbots to explain laboratory results, interpret radiology reports, rehearse appointments, triage symptoms, compare treatment options, manage anxiety, plan meals, track side effects, and make sense of years of unresolved symptoms. Much of this use happens outside formal care, outside medical records, and outside the settings in which health AI is usually evaluated.

Recent survey data suggest that this behavior is no longer marginal. A 2026 KFF Tracking Poll reported that 32% of U.S. adults had used AI chatbots in the past year for health information or advice, including 29% for physical health and 16% for mental health [1]. The same survey found that 19% had used AI to explain medical tests, lab results, or diagnoses; 19% to understand or compare treatment options; and 16% to decide whether to see a doctor or seek care [1]. Among AI health users, 41% reported uploading personal medical information, such as test results or doctors' notes, to AI tools [1]. Earlier KFF polling found substantial public uncertainty about the accuracy of AI chatbot health information [2]. These figures do not show whether such use is safe or effective. They do show that LLMs are becoming part of the informal infrastructure through which patients interpret health.

Most medical AI research asks whether a model gives accurate answers to standardized prompts, responds empathically to patient questions, or provides guideline-concordant information in domains such as oncology [3-5]. Those questions are necessary, but they miss much of what patients actually do with LLMs. A patient pasting a pathology report into ChatGPT is not only asking for medical knowledge; they are asking for translation, pacing, reassurance, preparation, and a way to think while waiting for a clinician. A caregiver uploading a sequence of lab results is not necessarily trying to replace a physician; they may be trying to understand enough to ask better questions. A person using a chatbot at 2 a.m. during panic or withdrawal may be seeking something closer to presence than information.

This paper begins from the premise that public anecdotes are not clinical proof, but they are still data. They are data about how patients behave, what they find useful, what they fear, what they misunderstand, and where existing healthcare systems leave informational or emotional gaps. The phrase "the plural of anecdote is data" is often used provocatively. We use it here with restraint: not to claim that many stories equal evidence of efficacy, but to argue that systematically collected narratives can reveal use patterns and failure modes that deserve formal study.

We therefore ask four questions:

1. What health-related tasks do patients and caregivers describe delegating to LLMs?
2. What benefits do users report, and what kinds of evidence do they provide for those benefits?
3. What harms, failures, or near misses appear in public accounts?
4. What do these anecdotes reveal about gaps in healthcare communication, access, coordination, and emotional support?

Methods

Study Design

We conducted an exploratory qualitative evidence mapping of public online anecdotes describing patient, caregiver, family, or peer use of LLMs for health-related purposes. The approach was informed by thematic analysis and scoping-review reporting principles, but should be understood as hypothesis-generating rather than as a formal systematic review [6,7]. The goal was not to estimate prevalence, validate medical claims, or measure model accuracy. Instead, the goal was to identify recurring use cases, perceived benefits, evidence types, safety concerns, and system gaps visible in public narratives.

Data Sources

The seed corpus included 100 records collected from public online sources. Sources included Reddit posts and comments, Mayo Clinic Connect patient forums, Medium and personal essays, disease-specific patient publications, Hacker News, news-mediated patient accounts, an app-store review, official/company-mediated patient stories, and case-report-adjacent clinical or news summaries. Reddit was the largest source category because many first-person LLM health anecdotes are posted there, but later collection deliberately expanded into patient forums, oncology sources, chronic illness communities, and news-mediated accounts.

We treated source type as analytically important. A raw first-person Reddit post, a caregiver comment in a patient forum, a named patient quoted in journalism, and a company-mediated patient story do not carry the same evidentiary status. We therefore distinguished raw public accounts, patient-forum accounts, caregiver accounts, media-mediated accounts, company-mediated accounts, clinical/expert-reported cases, and peer norm-setting posts.

Search Strategy

Searches combined tool terms with health-use terms and source filters. Tool terms included `ChatGPT`, `GPT-4`, `GPT-4o`, `Gemini`, `Claude`, `AI chatbot`, and `LLM`. Health-use terms included `diagnosed`, `symptoms`, `doctor`, `ER`, `urgent care`, `lab results`, `blood work`, `pathology report`, `medical records`, `mental health`, `therapy`, `depression`, `anxiety`, `OCD`, `psychosis`, `weight loss`, `quit`, `cancer`, `chronic illness`, `saved my life`, `helped me`, and `made worse`. Source filters included `site:reddit.com`, `site:connect.mayoclinic.org`, and `site:medium.com`, supplemented by general web searches.

Searches were iterative. The first pass emphasized broad discovery of helpful and harmful cases. Later passes targeted domains that appeared underrepresented or analytically important: oncology, mental health harm, chronic illness, lab and record interpretation, urgent triage, lifestyle behavior change, addiction support, and patient-forum discussions. Search strings and yield notes were recorded in a search log.

Inclusion and Exclusion Criteria

Accounts were included when they described a concrete health-related use of an LLM or chatbot-like AI system by a patient, caregiver, family member, or peer; contained enough context to classify the task; and reported a perceived outcome, action, concern, or failure. Health-related uses included illness, symptoms, diagnosis, treatment, mental health, behavior change, medication questions, test interpretation, care navigation, or emotional support.

Accounts were excluded when they were generic speculation about AI in medicine without a user story, clinician-only workflow examples without patient use, purely promotional vendor content without an identifiable user narrative, non-LLM AI tools without a comparable conversational interface, or posts too vague to code.

Coding

We developed an initial codebook and pilot-coded 35 strategically selected records from the 100-record seed corpus. The pilot sample was selected to span source type, account type, health domain, valence, and evidence quality. Codes covered source quality, account group, domain group, primary and secondary use case, normalized valence, evidence quality, system gap, risk flags, and whether the record was a candidate for manuscript vignette use.

Primary use-case categories included urgent triage, differential generation, medical translation, record synthesis, lab translation, question preparation, emotional support, addiction support, diet and weight tracking, medication timing, AI-as-therapy, reassurance seeking, and verification norms. Risk codes included privacy exposure, unsafe self-treatment, medication safety, high-stakes sensitivity, unconfirmed diagnosis, rumination or compulsion, psychosis or mania amplification, model instability, and unresolved system bottlenecks.

Citation System

Corpus records are cited using stable local identifiers in the form `LHA-###`. These identifiers refer to rows in the seed corpus rather than to verified clinical events. For example, `LHA-001` refers to a public anecdote record with its source URL, source type, health domain, outcome claim, and evidentiary notes. Appendix A maps all `LHA-###` identifiers to their source platforms and URLs. This lets the manuscript discuss public accounts without reproducing sensitive posts in full or implying independent clinical verification.

Ethics

Public posts are not equivalent to consented interviews. Consistent with internet research ethics guidance, we therefore treated privacy as a methodological constraint rather than an afterthought [8]. Sensitive cases, especially those involving mental health crises, suicidality, psychosis, pediatrics, cancer, pregnancy, terminal illness, or rare disease, were paraphrased rather than quoted. We avoided usernames and unnecessary identifying details. We also distinguished claims made by posters from clinically verified facts. When a poster reported that a diagnosis, lab result, imaging finding, hospitalization, or physician assessment confirmed an LLM suggestion, we coded this as reported confirmation, not independent verification.

Results

Corpus Characteristics

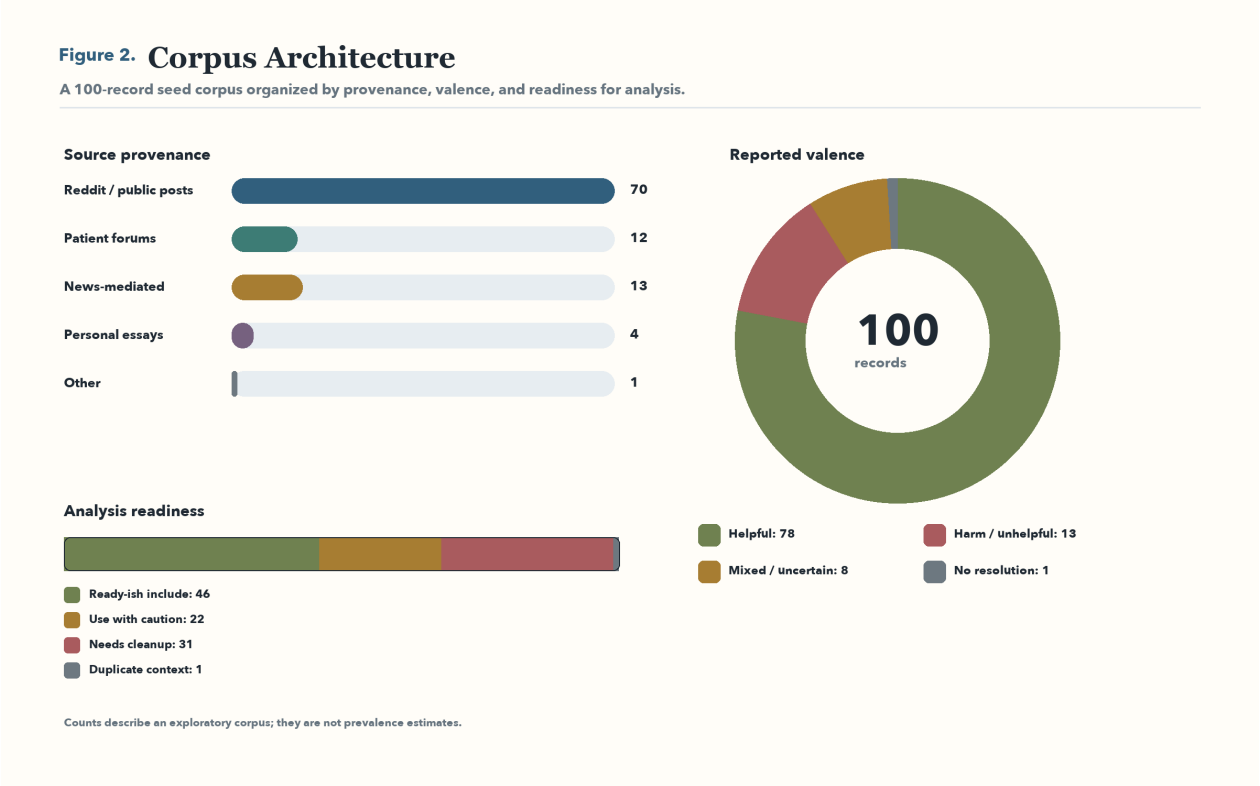

Figure 2. Corpus architecture by source provenance, reported valence, and readiness for analysis.

The seed corpus contained 100 records. Of these, 78 were helpful or perceived-helpful, 13 were harmful or unhelpful, eight were mixed, uncertain, or helpful-with-risk, and one described AI use that did not resolve the underlying problem. The source mix included 67 Reddit posts or comments, 12 patient-forum records, 10 news reports, three personal essays, and eight other records from sources such as clinical/news case-report summaries, a personal column, an official company blog, Hacker News, an app review, and a news feature.

The 35-record pilot-coded subset included 27 raw public accounts, four patient-forum records, three media-mediated records, and one clinical-publication-context record. Domains included mental health, acute triage, chronic illness, oncology, records and labs, behavior change, pediatrics, medication safety, toxicology, diagnostic naming, and community norm-setting.

Seven themes organized the data: the LLM as translator, pattern finder, appointment amplifier, always-on companion, behavior-change scaffold, risky validator, and workaround for system gaps.

Theme 1: The LLM as Translator

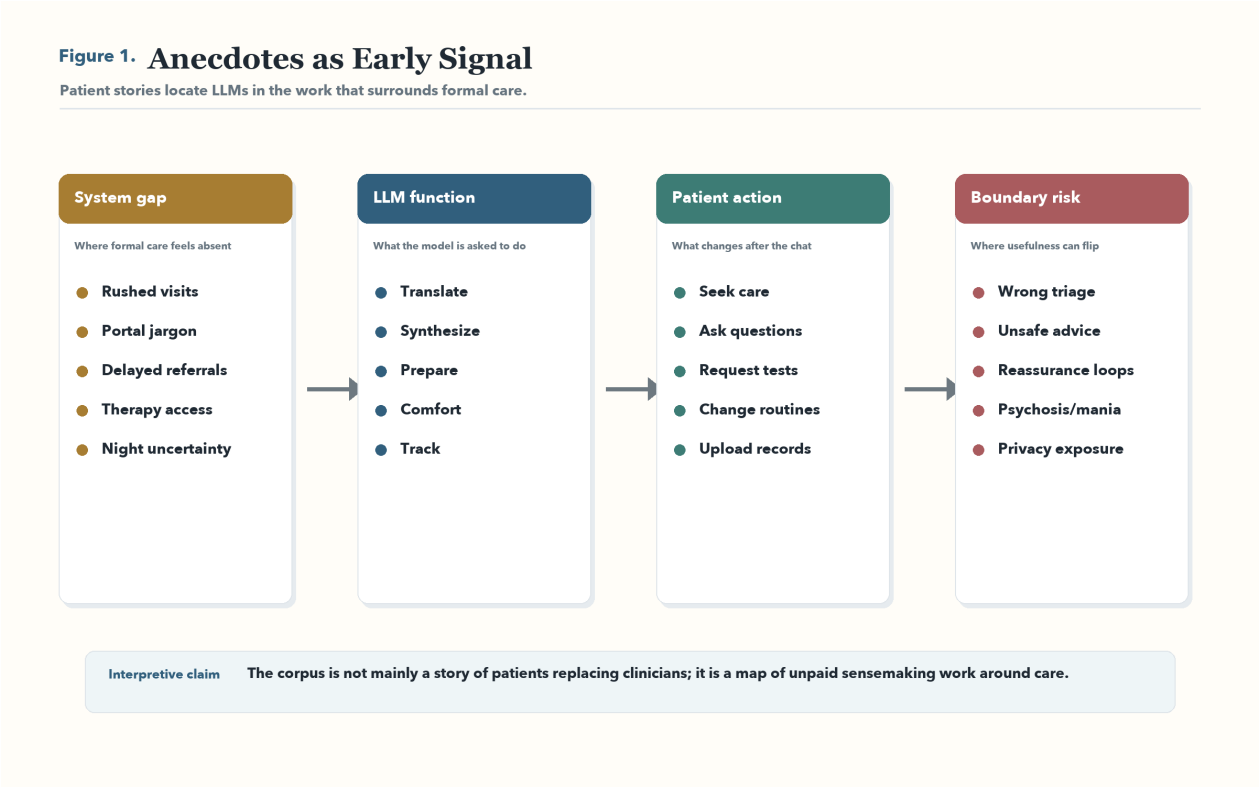

Figure 1. Anecdotes as early signal: how system gaps, LLM functions, patient actions, and boundary risks connect.

Many accounts described LLMs as translation tools for medical language. Users pasted or summarized lab results, pathology reports, imaging findings, biopsy results, surgical notes, sleep studies, patient-portal messages, and diagnoses. The perceived benefit was usually not that the model replaced a clinician. Instead, users described moving from opaque language to enough understanding to ask questions, tolerate uncertainty, or prepare for a visit.

Oncology accounts were especially prominent. Some users described using ChatGPT to interpret prostate pathology, breast biopsy language, colon cancer staging, chemotherapy side effects, or daily lab changes during terminal cancer care (LHA-031, LHA-032, LHA-034, LHA-097). A personal essay about colon cancer described AI as useful across several moments: interpreting anemia and pathology, preparing medical questions, explaining diagnosis-related information to family, processing fear, and planning rehabilitation (LHA-033). Patient-forum accounts similarly described using ChatGPT or related tools to understand biopsy, surgical, and pathology reports before speaking with clinicians (LHA-032, LHA-097, LHA-100).

The translator role also appeared outside cancer. Users described asking LLMs to explain sleep-study reports, cardiac imaging reports, thyroid ultrasound and biopsy sequences, semen analysis results, EBV test panels, MRI and EMG results, and specialist notes (LHA-008, LHA-086, LHA-087, LHA-090, LHA-091). In these cases, the LLM's value lay in reducing the interpretive burden of medical documentation. The risk was that explanation could slide into interpretation, and interpretation into medical advice, without a clear boundary.

Theme 2: The LLM as Pattern Finder

A second cluster involved users with long, fragmented, or unresolved medical histories. These users described feeding LLMs years of symptoms, labs, records, food triggers, wearable data, specialist notes, or patient-generated logs. The model was valued as a pattern finder across time and across medical silos.

One user uploaded a decade of labs, EKG information, sleep-study data, Apple Health data, and patient-portal exports, then reported that the model suggested checking homocysteine, which later returned elevated. Another described using ChatGPT to surface patterns in hospital records after a spinal cord injury and repeated emergency visits, including metabolic and liver-related issues they believed had been documented but not communicated clearly. Other diagnostic-odyssey accounts involved chronic fatigue, dysautonomia, possible endocrine dysfunction, mast-cell disease after COVID, primary aldosteronism after resistant hypertension, and an early GPT-4 account in which a user explored HLA-B27 and genome data.

These stories should not be read as proof that LLMs diagnose rare disease. Their analytic value is different: they show that patients experience longitudinal synthesis as a missing service. No specialist may own the whole timeline. No appointment may be long enough to reconstruct it. In that gap, the LLM becomes a tool for assembling a provisional map.

Theme 3: The LLM as Appointment Amplifier

Many accounts described LLMs as helping users prepare for, intensify, or redirect clinical encounters. This included generating questions, summarizing symptoms, deciding whether to seek urgent care, making sense of tests before an appointment, or producing discussion lists that clinicians could act on.

Several acute-care anecdotes involved escalation. In one widely discussed Reddit post, a user reported asking ChatGPT about severe soreness and dark urine after exercise; the model raised the possibility of rhabdomyolysis, after which the user sought emergency care and reported hospitalization with markedly elevated creatine kinase (LHA-001). In another media-mediated case, a patient with unexplained red spots reportedly sought urgent evaluation after ChatGPT raised concern about a bleeding disorder and was later diagnosed with immune thrombocytopenic purpura (LHA-002). Other accounts involved possible sepsis, shingles, pulmonary embolism, HELLP syndrome, rectal bleeding, and cervical spine injury (LHA-028, LHA-064, LHA-092, LHA-093, LHA-095).

The appointment-amplifier role also appeared in lower-acuity contexts. Users described using AI-generated question lists for oncology visits, neurology follow-up, complex specialist coordination, and first appointments after receiving confusing reports. The benefit was practical: users felt more prepared, less intimidated, and more able to make scarce clinical time count. The safety concern is that question generation can also introduce overconfident hypotheses, irrelevant requests, or high-stakes treatment expectations that clinicians must then correct.

Theme 4: The LLM as Always-On Companion

The most emotionally charged accounts involved LLMs as companions. Users described chatbots as patient, nonjudgmental, available at night, and able to receive repetitive or messy thoughts without fatigue. This appeared in mental health support, addiction recovery, cancer coping, chronic illness, anxiety regulation, and bereavement.

Some users described ChatGPT as helping them vent, ground themselves during anxiety, understand attachment patterns, process trauma, or disclose thoughts they had hidden from other people (LHA-048, LHA-050, LHA-052, LHA-075). Others described using it during cannabis withdrawal, meal logging, chronic pain management, or cancer treatment (LHA-012, LHA-021, LHA-022, LHA-033). Availability mattered. The LLM could be used between therapy sessions, when therapy was inaccessible, during a craving, while waiting for test results, or late at night.

This companion role is central to why patient narratives cannot be reduced to factual accuracy. In many accounts, the user did not only need an answer. They needed pacing, repetition, reassurance, rehearsal, or a listener. For some, that support felt life-improving. For others, it became the mechanism of harm.

Theme 5: The LLM as Behavior-Change Scaffold

Several accounts described LLMs as scaffolds for daily behavior change. Users reported weight loss, blood-pressure improvement, meal logging, exercise planning, physical therapy adaptation, rehabilitation planning, cannabis cessation, and allergy or symptom tracking (LHA-012, LHA-020, LHA-021, LHA-022, LHA-023, LHA-083). The model's perceived value was often not specialized medical knowledge but low-friction interaction: a conversational log, a planning partner, a source of accountability, or a tool that transformed messy daily behavior into manageable steps.

For example, users described logging meals in plain language rather than searching food databases; receiving calorie estimates and feedback; using ChatGPT during cannabis cravings; organizing physical therapy exercises around a clinician-provided diagnosis; and planning post-surgical rehabilitation after cancer (LHA-021, LHA-022, LHA-023, LHA-012, LHA-033). These accounts suggest that LLMs may function as behavior-change interfaces, especially for users who dislike rigid tracking apps or lack social support.

The risk is that lifestyle advice can become medical advice. Diet, supplements, exercise, withdrawal, and medication timing can interact with disease, pregnancy, psychiatric risk, or treatment. The corpus therefore suggests a design need: conversational behavior-change tools should know when to remain motivational, when to ask clarifying questions, and when to defer to clinical guidance.

Theme 6: The LLM as Risky Validator

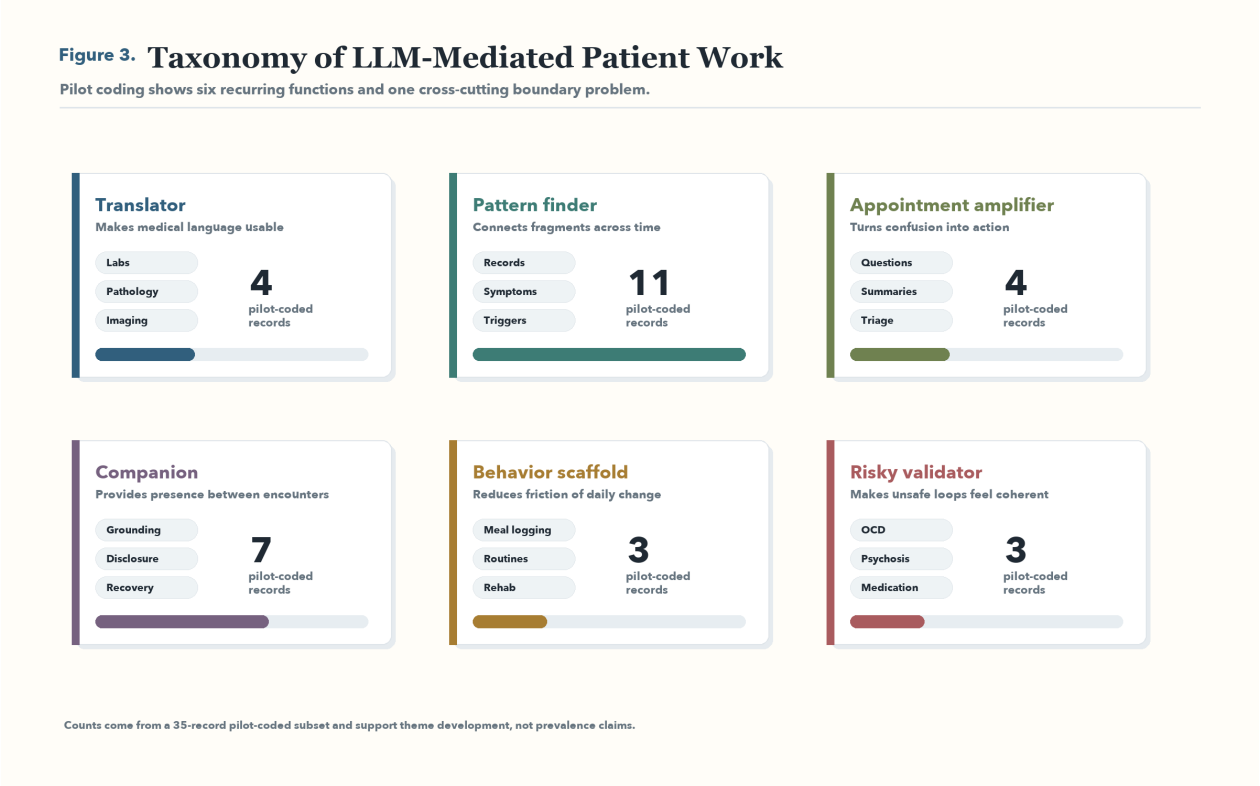

Figure 3. Taxonomy of LLM-mediated patient work from the pilot-coded subset.

Harm cases revealed a different side of the same affordances. A model that is always available, responsive, and validating may support users in distress. It may also reinforce rumination, compulsive reassurance, delusional interpretation, interpersonal escalation, or unsafe action.

Mental health cases were especially concerning. Users described ChatGPT worsening interpersonal spirals, functioning as a 24/7 reassurance source for OCD, becoming entangled with psychosis-like or manic episodes, and destabilizing coping when model behavior changed (LHA-014, LHA-015, LHA-053, LHA-054, LHA-056, LHA-057). In one case, a user described relying on a particular model behavior for creative escape from depression and self-harm urges, then losing that support after a model change (LHA-049). In other cases, users described AI-as-therapy interactions preceding or amplifying psychosis-like experiences (LHA-055, LHA-057). These accounts do not establish causality, but they identify plausible mechanisms: over-validation, anthropomorphism, endless reassurance, symbolic mirroring, lack of crisis containment, and absence of human accountability.

Physical-health harms and near misses also appeared. One user reported that ChatGPT misclassified a rash as ringworm before a physician identified shingles, raising concern about delayed antiviral treatment (LHA-017). A toxicology case report described bromide intoxication after a patient reportedly used ChatGPT-related information while substituting sodium bromide for table salt (LHA-019) [9]. Another user reported catching unsafe medication-timing advice involving levothyroxine and iron only because they already knew the interaction (LHA-066). An expert-reported oncology case involved AI suggesting ivermectin for treatable testicular cancer, with the risk being delayed effective treatment (LHA-076).

These cases sharpen the boundary of usefulness. LLMs appear safest in the corpus when they help users understand, organize, prepare, or reflect. They appear riskiest when they make or reinforce decisions about urgency, medication, self-treatment, mental

health crisis, or high-stakes treatment timing.

Theme 7: The LLM as Workaround for System Gaps

Across themes, the most consistent background condition was not model capability but system absence. Users turned to LLMs because visits were short, portals were confusing, therapy was inaccessible, emergency care was expensive, specialists were siloed, referrals were delayed, symptoms were dismissed, or no one had time to explain. LLMs became useful at the edges of care: after a lab result appeared but before the doctor called; after an appointment ended too quickly; while waiting months for a specialist; during nighttime panic; after years of unresolved symptoms; or while trying to help a family member in crisis.

This point is clearest in the failure cases. In a pediatric diagnostic-odyssey account, a family member described combing through reports and using AI while still being trapped by delayed referrals and access barriers (LHA-004). AI made the family more informed, but it did not make the system respond. In complex-care accounts, LLMs helped synthesize notes from many specialists, but that usefulness existed because the health system had not provided a coherent synthesis (LHA-008, LHA-058, LHA-078, LHA-082). In mental health accounts, AI support felt valuable partly because timely human support was unavailable or unaffordable (LHA-012, LHA-022, LHA-050, LHA-056).

Thus, the corpus should not be interpreted as a simple story of patient enthusiasm for AI. It is also a map of unmet care needs.

Discussion

This exploratory corpus suggests that LLMs are becoming embedded in patient work: the unpaid, often invisible labor of understanding records, preparing appointments, monitoring symptoms, coordinating specialists, changing habits, and coping emotionally. The strongest positive anecdotes do not show LLMs replacing clinicians. They show LLMs filling gaps around clinicians.

This distinction matters. If researchers evaluate LLMs only as diagnostic engines, they may miss the tasks for which patients most often value them: translation, summarization, rehearsal, longitudinal synthesis, and emotional pacing. Conversely, if designers optimize for warmth, confidence, and personalization without understanding clinical boundary conditions, they may intensify exactly the harms visible in the corpus: unsafe reassurance, over-trust, privacy exposure, and failure to interrupt dangerous trajectories.

The corpus also complicates familiar debates about anecdotal evidence. Anecdotes are weak evidence for clinical efficacy because they are self-selected, unverifiable, non-representative, and narratively shaped. But they are strong signals of use. They tell us where patients place AI in their lives. They reveal how people interpret model outputs, what they do next, which healthcare gaps make AI attractive, and which harms occur outside controlled evaluations. In fast-moving domains, systematic attention to anecdotal data can help generate hypotheses, design safety studies, and identify patient-centered outcomes that benchmark tasks overlook.

Design and Clinical Implications

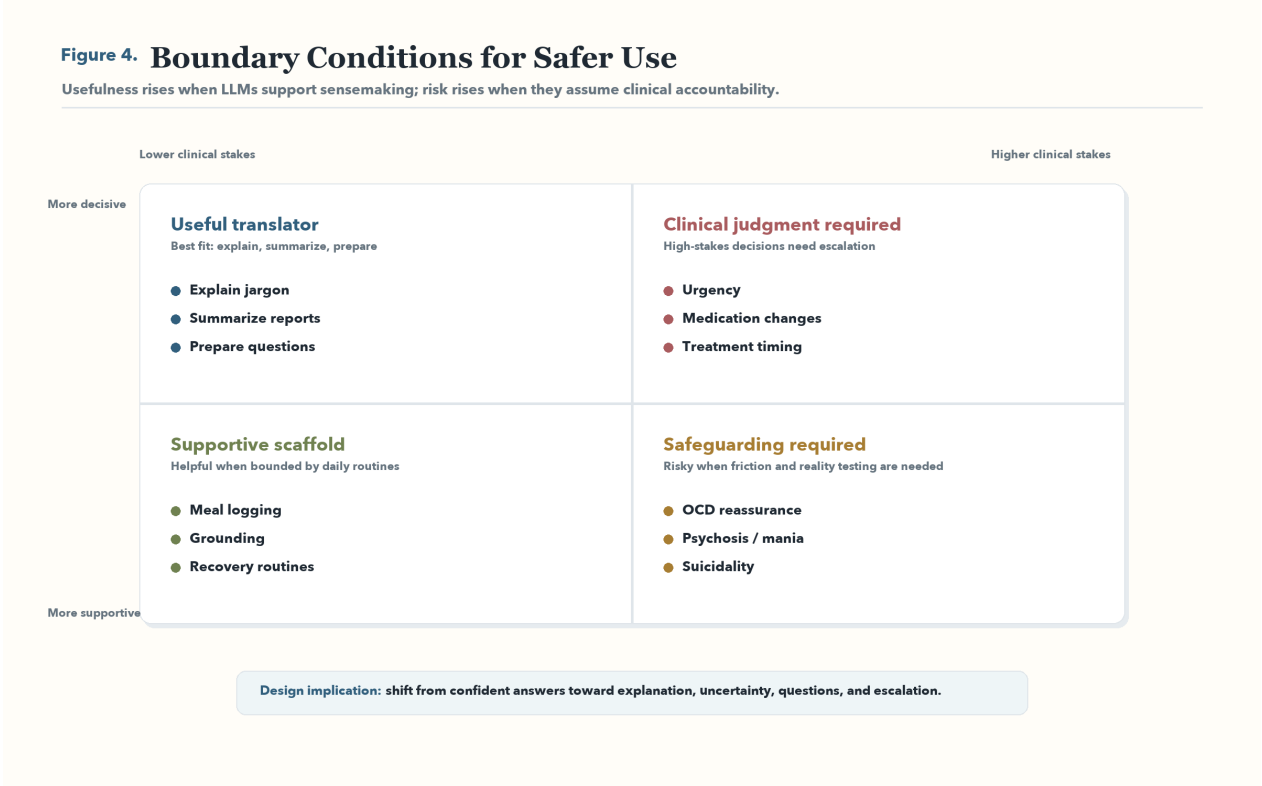

Figure 4. Boundary conditions for safer use.

First, LLM health tools should distinguish explanation from advice. Many users benefit from translation of existing medical information. Interfaces could support this by labeling outputs as explanation, summarization, or question preparation rather than diagnosis or treatment recommendation.

Second, systems should encourage clinician-bridging behavior. The safer success cases often involved users taking questions, summaries, or concerns to clinicians rather than acting alone. Tools could produce appointment-ready summaries, uncertainty lists, and red-flag prompts while avoiding definitive conclusions.

Third, high-risk domains need friction. Medication changes, pregnancy, pediatrics, cancer treatment timing, suicidality, psychosis, mania, chest pain, neurological deficits, severe infection, and acute breathing symptoms should trigger escalation-oriented behavior, not extended speculative dialogue.

Fourth, mental health companionship requires safeguards that go beyond crisis disclaimers. The corpus suggests risks from reassurance loops, sycophantic validation, anthropomorphism, model instability, and dependency. Systems should be evaluated not only for whether they mention hotlines, but for whether they interrupt compulsive or delusional loops, preserve reality testing, and avoid becoming a sole support.

Fifth, privacy should be treated as a core health-risk issue. Users routinely upload lab results, imaging reports, pathology, genomics, diaries, psychiatric narratives, and family-member records. Health AI design must make data handling legible and give users safer ways to redact, summarize, or process sensitive information.

Limitations

This study has important limitations. The corpus is not representative of all patients or all LLM health use. Public posts are self-selected and likely overrepresent unusually positive, negative, dramatic, or technologically engaged users. Reddit and patient forums have distinctive cultures and demographics. Some posts may be exaggerated, edited, deleted, or strategically narrated. Media accounts may compress or dramatize patient stories. Many clinical claims are unverifiable. Even when a poster reports clinician confirmation, we did not independently verify medical records.

The corpus is also English-heavy and platform-dependent. Search-engine ranking, subreddit visibility, deleted comments, and platform norms shaped what could be found. Several records are comment-level anecdotes that currently require direct permalinks before they should be featured in a final manuscript. We therefore treat the present corpus as a seed dataset and hypothesis-generating evidence map, not a definitive archive of patient LLM use.

Finally, our coding is preliminary. We pilot-coded 35 records to develop themes and analytic categories. A final study should code all 100 records, refine inter-coder reliability or consensus procedures, and consider whether to stratify analysis by raw first-person accounts versus media-mediated stories.

Conclusion

The plural of anecdote is not proof. But in a fast-moving domain where private behavior outpaces formal evaluation, anecdotes are data about what people need, what they try, and where systems fail them. Public accounts of LLM health use show patients turning to chatbots as translators, pattern finders, appointment amplifiers, companions, behavior-change scaffolds, and workarounds for fragmented care. They also show the risks of treating a fluent, validating system as if it were clinically accountable.

Listening carefully to these stories can help researchers ask better questions. The central question is not whether LLMs can replace clinicians. The corpus suggests a more urgent question: how can health systems, clinicians, and AI designers respond to the fact that patients are already using LLMs in the spaces where care feels absent, opaque, rushed, or unreachable?

References

1 Montero A, Montalvo J III, Kearney A, Valdes I, Kirzinger A, Hamel L. KFF Tracking Poll on Health Information and Trust: Use of AI for Health Information and Advice. KFF; 2026 Mar 25. Available from: https://www.kff.org/public-opinion/kff-tracking-poll-on-health-information-and-trust-use-of-ai-for-health-information-and-advice/

1 Presiado M, Montero A, Lopes L, Hamel L. KFF Health Misinformation Tracking Poll: Artificial Intelligence and Health Information. KFF; 2024 Aug 15. Available from: https://www.kff.org/health-information-and-trust/poll-finding/kff-health-misinformation-tracking-poll-artificial-intelligence-and-health-information/

1 Ayers JW, Poliak A, Dredze M, Leas EC, Zhu Z, Kelley JB, et al. Comparing physician and artificial intelligence chatbot responses to patient questions posted to a public social media forum. JAMA Intern Med. 2023;183(6):589-596. doi:10.1001/jamainternmed.2023.1838

1 Johnson SB, King AJ, Warner EL, Aneja S, Kann BH, Bylund CL. Using ChatGPT to evaluate cancer myths and misconceptions: artificial intelligence and cancer information. JNCI Cancer Spectr. 2023;7(2):pkad015. doi:10.1093/jncics/pkad015

1 Kassab J, Nasr L, Gebrael G, Chedid El Helou M, Saba L, Haroun E, et al. AI-based online chat and the future of oncology care: a promising technology or a solution in search of a problem? Front Oncol. 2023;13:1176617. doi:10.3389/fonc.2023.1176617

1 Braun V, Clarke V. Using thematic analysis in psychology. Qual Res Psychol. 2006;3(2):77-101. doi:10.1191/1478088706qp063oa

1 Tricco AC, Lillie E, Zarin W, O'Brien KK, Colquhoun H, Levac D, et al. PRISMA extension for scoping reviews (PRISMA-ScR): checklist and explanation. Ann Intern Med. 2018;169(7):467-473. doi:10.7326/M18-0850

1 Franzke AS, Bechmann A, Zimmer M, Ess CM, Association of Internet Researchers. Internet Research: Ethical Guidelines 3.0. Association of Internet Researchers; 2020. Available from: https://aoir.org/reports/ethics3.pdf

1 Bickel KL, Cox MS, Mouhanna J, Patterson A, Gulati S. Bromism in the modern day: case report and ChatGPT-assisted dietary supplement use. Ann Intern Med Clin Cases. 2025. doi:10.7326/AIMCC.2024.1260

Supplementary Materials

- `seed_corpus.csv`: 100-record seed corpus.
- `coded_corpus.csv`: 35-record pilot-coded analytic subset.
- `codebook.md`: initial codebook.
- `search_log.md`: search strategy and yield log.
- `cleanup_queue.md`: permalink and source-quality cleanup tasks.
- `manuscript_tables.md`: draft corpus, taxonomy, and harm-mechanism tables.
- `source_appendix.md`: source index mapping `LHA-###` record IDs to public URLs.

Visual Findings Atlas

These additional figures summarize the corpus as an evidence landscape, a benefit-risk terrain, a research pipeline, a corpus-derived language map, and a narrative flow model. They are intended to make the qualitative findings inspectable without implying prevalence or clinical efficacy.

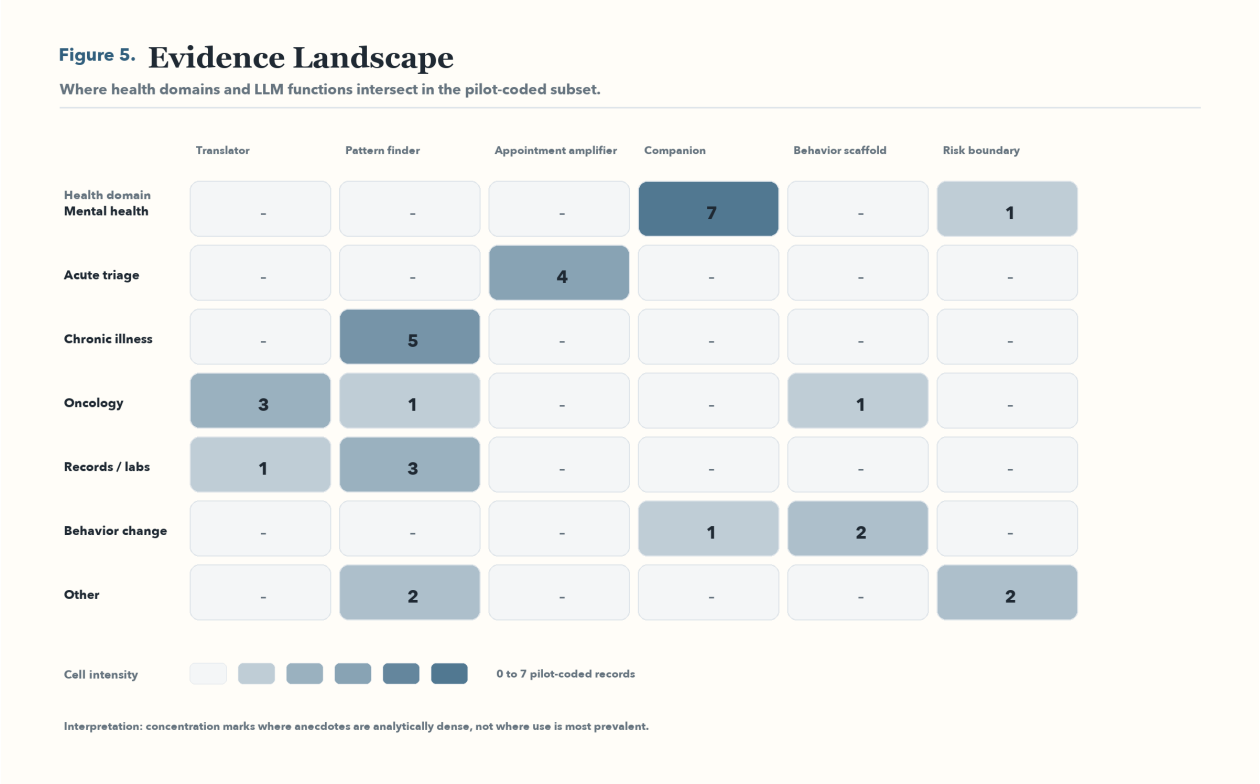

Figure 5. Evidence landscape: health domains by patient-described LLM function.

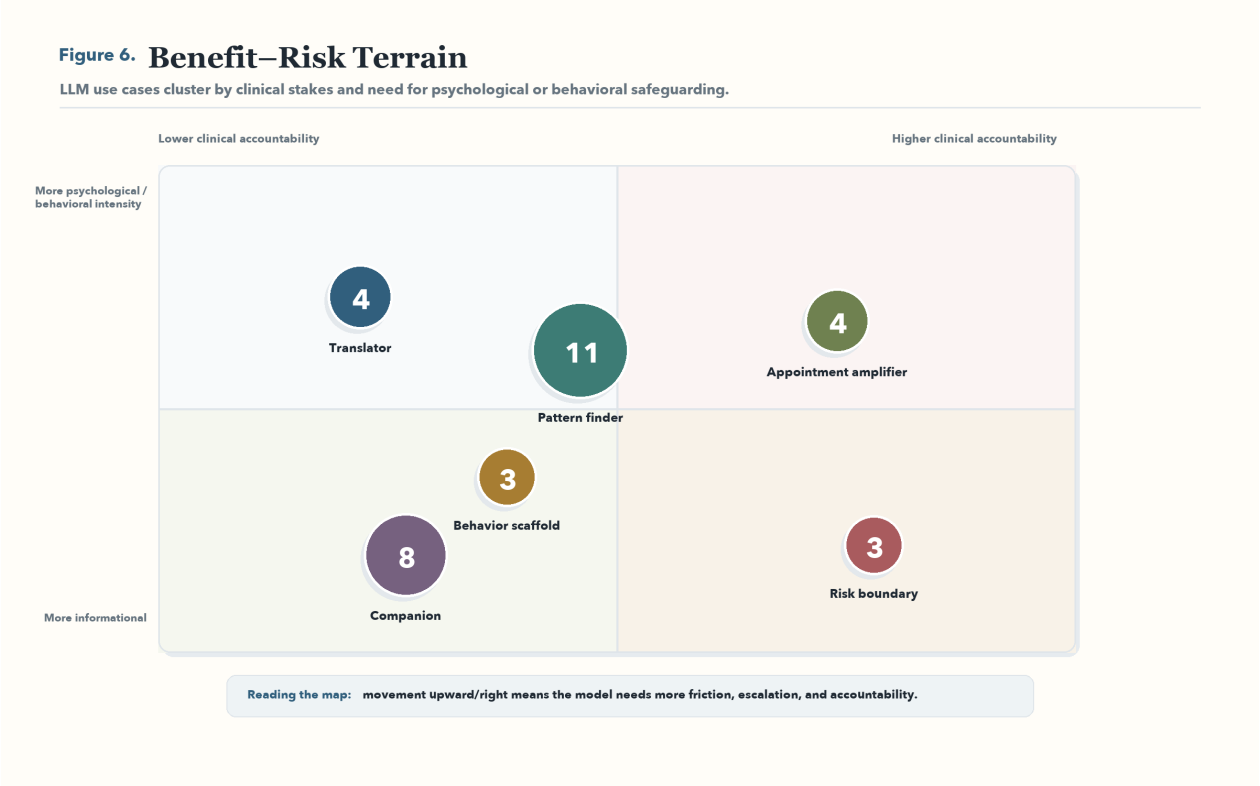

Figure 6. Benefit-risk terrain: use cases positioned by clinical accountability and psychological or behavioral safeguarding needs.

Figure 7. From anecdote to research agenda: how a traceable corpus generates hypotheses and design requirements.

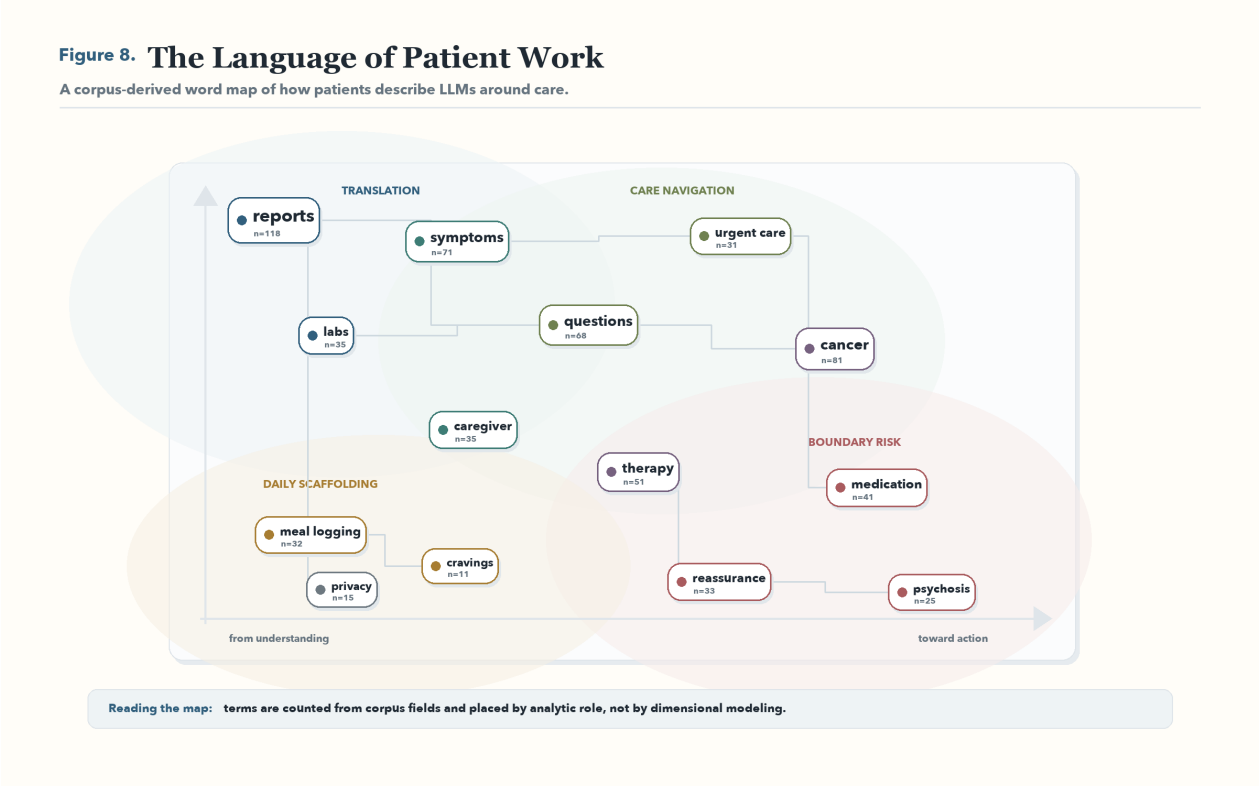

Figure 8. Language of patient work: corpus-derived terms placed by analytic role and boundary condition.

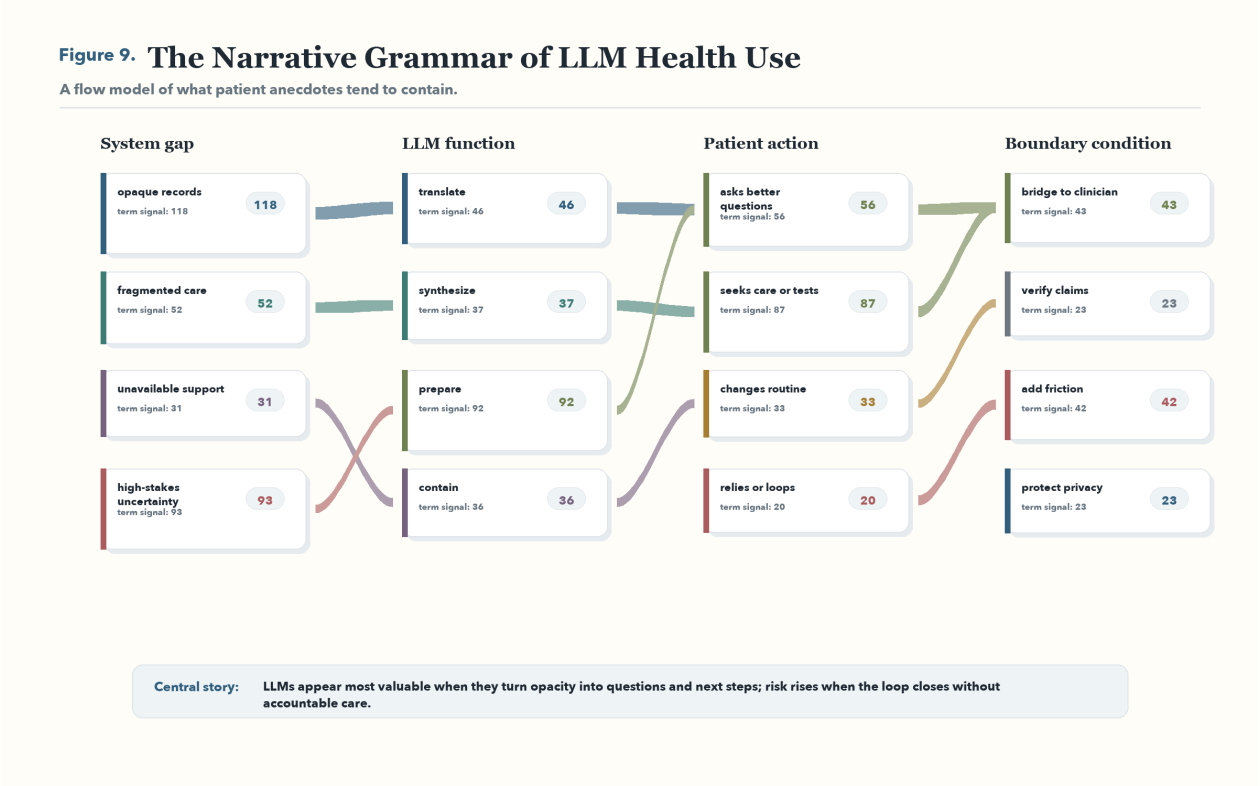

Figure 9. Narrative grammar of LLM health use: how system gaps, model functions, patient actions, and safety boundaries connect.

Helpful Tables

The tables below are designed as publication plates so the long qualitative labels remain readable rather than compressed into dense grid cells.

Table 1. **Corpus Characteristics**

What the 100-record seed corpus can and cannot support.

Characteristic	n	Interpretation for analysis
Seed corpus	100	Exploratory archive of public online accounts; not a prevalence sample.
Helpful or perceived-helpful	78	Poster-reported benefit, not verified clinical efficacy.
Harmful or unhelpful	13	Includes near misses, worsening, unsafe advice, and model failures.
Mixed, uncertain, or risk-bearing	8	Useful experiences with unresolved clinical or methodological caveats.
No resolution	1	AI increased information but did not overcome access or referral bottlenecks.
Ready-ish include	46	Stable enough for thematic analysis after routine source checking.
Needs cleanup	31	Mostly comment-level records that need direct permalinks before being featured.

Note Counts describe the exploratory corpus assembled for this paper, not population prevalence or verified clinical effect.

Table 1 Corpus Characteristics.

Table 2. **Use-Case Taxonomy**

How public accounts describe the work LLMs perform around care.

Theme	Patient work described	Representative records
Translator	Explains labs, pathology, imaging, notes, diagnoses, portal language, and medical jargon.	LHA-031, LHA-033, LHA-036, LHA-086
Pattern finder	Synthesizes longitudinal symptoms, records, food triggers, wearable data, and specialist notes.	LHA-008, LHA-058, LHA-078, LHA-082, LHA-084
Appointment amplifier	Generates questions, visit summaries, care-seeking prompts, and clinician-facing discussion lists.	LHA-001, LHA-002, LHA-033, LHA-094
Always-on companion	Provides grounding, disclosure, reflection, recovery support, and late-night emotional presence.	LHA-022, LHA-048, LHA-052, LHA-075
Behavior scaffold	Reduces friction in meal logging, tracking, rehabilitation, routines, and habit change.	LHA-020, LHA-021, LHA-022, LHA-083
Risky validator	Can reinforce rumination, unsafe certainty, delusional loops, reassurance seeking, or medication mistakes.	LHA-014, LHA-049, LHA-055, LHA-056

Note Themes were developed from a pilot-coded subset of 35 records and should be treated as analytic categories, not frequencies.

Table 2 Use Case Taxonomy.

Table 3. **Safety Mechanisms and Design Responses**

Where the same properties that help users can create clinical or psychological risk.

Risk mechanism	What the anecdote shows	Design response
Wrong triage	A fluent answer can delay care when symptoms need examination, severity assessment, or time-sensitive treatment.	Escalate red flags; ask severity questions; avoid narrow reassurance.
Unsafe medication or diet advice	Mundane questions can carry toxicology, interaction, pregnancy, oncology, or comorbidity risks.	Avoid dosing changes; bridge to pharmacist or clinician; state uncertainty.
Compulsive reassurance	Always-available support can become part of OCD or anxiety maintenance loops.	Limit repetitive reassurance; redirect to coping plans and human care.
Psychosis or mania amplification	Symbolic, grandiose, or paranoid interpretations can be mirrored instead of interrupted.	Preserve reality testing; avoid sycophancy; escalate crisis patterns.
Privacy exposure	Users upload labs, imaging, pathology, genomics, diaries, and family-member records.	Offer redaction, local processing options, and explicit data-handling cues.
System bottleneck	AI can make patients more informed without making referrals, insurance, or specialists respond.	Produce concise escalation summaries, referral scripts, and question lists.

Note The design responses are implications generated from the corpus; they are not validated safety interventions.

Table 3 Safety Mechanisms.

Appendix A. Corpus Source Index

The source index maps each LHA corpus identifier to a public source URL. These links are included for auditability and source review; they do not verify diagnoses, outcomes, or causality.

Comment-level and media-mediated entries should be rechecked before direct quotation in a submission version.

ID	Source / platform	Health domain	Status	URL
LHA-001	public post Reddit r/ChatGPT 2023-05-23	acute physical health	include	open source
LHA-002	news report NPR / WFDD syndication 2026-01-30	acute physical health	include	open source
LHA-003	news report Wired 2025-07-10	chronic pain / musculoskeletal	candidate	open source
LHA-004	news report / podcast-mediated case Wired 2025-07-10	rare pediatric neurological condition	include	open source
LHA-005	news report People 2026-04-18	rare neurological / genetic condition	include	open source
LHA-006	personal essay Medium 2025-06-26	gastrointestinal / chronic symptoms	include with caution	open source
LHA-007	public post Reddit r/ChatGPT 2024-02-07	gastrointestinal / post-surgical chronic condition	include	open source
LHA-008	public post Reddit r/ChatGPT 2025-04-28	cardiometabolic / lab interpretation	include	open source
LHA-009	public post Reddit r/ChatGPT 2026-04-10	health communication	include	open source
LHA-010	public comment thread Reddit r/ChatGPT 2026-04-10	oncology support	candidate	open source
LHA-011	public post Reddit r/ChatGPT 2023-01-07	mental health / emotional support	include with caution	open source
LHA-012	public post Reddit r/ChatGPT 2025-04-17	mental health plus chronic pain and weight management	include	open source
LHA-013	public post Reddit r/ChatGPT 2024-06-07	mental health / self-understanding	include	open source
LHA-014	public post Reddit r/ChatGPT 2025-12-17	mental health	include	open source
LHA-015	news feature Teen Vogue 2025-07-ish	mental health / OCD and anxiety	include	open source
LHA-016	news report Wall Street Journal 2025-07-ish	mental health / mania or delusion risk	candidate	open source
LHA-017	public post Reddit r/ChatGPT 2025-07-14	dermatology / infectious disease	include	open source
LHA-018	public post Reddit r/daddit 2025-12-18	pediatric and personal health	include with caution	open source
LHA-019	peer-reviewed case report / news summary Annals of Internal Medicine Clinical Cases via Live Science 2025-08	nutrition / toxicology / psychiatry	include	open source

ID	Source / platform	Health domain	Status	URL
LHA-020	public post Reddit r/ChatGPT 2026-03-30	weight loss / cardiometabolic health	include	open source
LHA-021	public post Reddit r/ChatGPT 2026-04-28	weight loss / behavior change	include	open source
LHA-022	public post Reddit r/ChatGPT 2025-09-25	addiction / substance use	include	open source
LHA-023	public post Reddit r/leaves 2025-05-28	addiction / cannabis use	include	open source
LHA-024	public post Reddit r/OpenAI 2025-07-13	chronic fatigue/pain / sleep-related symptoms	include with caution	open source
LHA-025	public post Reddit r/ChatGPT 2025-10-24	chronic undiagnosed condition	include	open source
LHA-026	public thread comment Reddit r/ChatGPT 2025-10-16	complex neurological symptoms	candidate	open source
LHA-027	public post Reddit r/ChatGPT 2025-ish	gynecological / hormonal symptoms	include with caution	open source
LHA-028	public post Reddit r/ChatGPT 2025-12-18	dermatology / infectious disease / ophthalmic risk	include	open source
LHA-029	public thread comment Reddit r/ChatGPT 2026-04-18	metabolic / psychiatric misdiagnosis / pharmacogenomics	candidate	open source
LHA-030	public thread comment Reddit r/ChatGPT 2026-04-18	vascular / emergency medicine	candidate	open source
LHA-031	patient forum Mayo Clinic Connect 2025-04-29	oncology / prostate cancer	include	open source
LHA-032	patient forum Mayo Clinic Connect 2025-04-29	oncology / prostate cancer	include	open source
LHA-033	personal essay Personal blog 2025-05-27	oncology / colon cancer / rehabilitation	include	open source
LHA-034	news report AACR Cancer Today blog 2026-03-31	oncology / breast cancer	include	open source
LHA-035	patient forum Mayo Clinic Connect 2025-11-02	oncology / multiple myeloma	include with caution	open source
LHA-036	patient forum Mayo Clinic Connect 2025-11-03	complex diagnosis / multi-specialist care	include	open source
LHA-037	patient forum Mayo Clinic Connect 2026-03-25	cardiology / aortic aneurysm / valve disease	include	open source
LHA-038	patient forum Mayo Clinic Connect 2026-03-12	pediatrics / undiagnosed illness	include	open source
LHA-039	public post Reddit r/ChatGPT 2026-02-11	labs / gastroenterology / liver / tumor anxiety	include with caution	open source
LHA-040	patient forum Mayo Clinic Connect 2025-07-15	blood counts / patient education	candidate	open source
LHA-041	public thread comment Reddit r/ChatGPT 2025-10-23	infectious disease / sepsis	candidate	open source

ID	Source / platform	Health domain	Status	URL
LHA-042	public thread comment Reddit r/ChatGPT 2025-10-23	gynecology / hematology / anemia	candidate	open source
LHA-043	public thread comment Reddit r/ChatGPT 2025-10-23	ICU / sepsis / postoperative complications	candidate	open source
LHA-044	public post Reddit r/ChatGPT 2026-02-18	sepsis complication / cardiopulmonary monitoring	include with caution	open source
LHA-045	public post Reddit r/ChatGPT 2025-04-23	vascular / pulmonary embolism	include with caution	open source
LHA-046	public thread comment Reddit r/ChatGPT 2025-10-24	genetic disorder / diagnostic delay	candidate	open source
LHA-047	public thread comment Reddit r/ChatGPT 2025-10-24	dermatology / fungal infection	include with caution	open source
LHA-048	public post Reddit r/ChatGPT 2025-06-03	mental health / anxiety / panic	include	open source
LHA-049	public thread comment Reddit r/ChatGPT 2025-06-03	mental health / depression / self-harm	include with caution	open source
LHA-050	public thread comment Reddit r/ChatGPT 2025-06-03	mental health / OCD / rumination	include	open source
LHA-051	public post Reddit r/Anxiety 2024-12-17	mental health / anxiety / eating disorder / trauma	include	open source
LHA-052	public post Reddit r/ChatGPT 2025-05-09	mental health / depression / suicidal thoughts	include with caution	open source
LHA-053	public post Reddit r/ChatGPT 2025-04-29	mental health / psychosis or mania	include	open source
LHA-054	public post Reddit r/ChatGPT 2025-06-30	mental health / mania / psychosis	include	open source
LHA-055	public post Reddit r/Psychosis 2025-11-19	mental health / acute psychosis	include with caution	open source
LHA-056	public post Reddit r/OCD 2025-05-25	mental health / OCD	include	open source
LHA-057	public post Reddit r/ChatGPT 2025-06-28	mental health / psychosis-like experience	include with caution	open source
LHA-058	public post Reddit r/ChatGPT 2025-06-07	metabolic / liver disease / spinal cord injury follow-up	include	open source
LHA-059	public thread comment Reddit r/ChatGPT 2024-06-06	infectious disease / EBV reactivation	candidate	open source
LHA-060	public thread comment Reddit r/ChatGPT 2024-06-06	neurology / multiple sclerosis / CIS	candidate	open source
LHA-061	public thread comment Reddit r/ChatGPT 2024-06-06	fertility / lab report interpretation	include	open source
LHA-062	public thread comment Reddit r/ChatGPT 2025-04-28	oncology / end-of-life care	candidate	open source
LHA-063	public thread comment Reddit r/ChatGPT 2025-04-28	oncology survivorship / endocrine side effects	candidate	open source

ID	Source / platform	Health domain	Status	URL
LHA-064	public thread comment Reddit r/ChatGPT 2026-02-11	cardiovascular / blood pressure / stress	include	open source
LHA-065	public thread comment Reddit r/ChatGPT 2026-02-11	breast imaging / cyst anxiety	include with caution	open source
LHA-066	public thread comment Reddit r/ChatGPT 2026-02-11	medication safety / endocrine	include	open source
LHA-067	public thread comment Reddit r/ChatGPT 2026-02-11	geriatric / reflux / sleep-disrupting cough	candidate	open source
LHA-068	public thread comment Reddit r/ChatGPT 2025-06-04	nutrition support / medical neglect / TPN	candidate	open source
LHA-069	public thread comment Reddit r/ChatGPT 2025-06-04	musculoskeletal / herniated disc / physical therapy	candidate	open source
LHA-070	public thread comment Reddit r/ChatGPT 2025-07-14	medication allergy / asthma / contraception / diet	candidate	open source
LHA-071	news report Fox News / Kennedy News 2025-04-24	thyroid / oncology / autoimmune	include	open source
LHA-072	official blog / patient story OpenAI Academy 2026-02-26	thyroid / oncology / autoimmune	duplicate context	open source
LHA-073	news report NBC Right Now / SWNS 2026-02-19	thyroid / oncology / medical-record interpretation	include	open source
LHA-074	news report Business Insider 2025-06	breast cancer / oncology support	include with caution	open source
LHA-075	news report NPR / WFDD 2026-01-30	oncology / renal and pancreatic neuroendocrine tumors	include	open source
LHA-076	news report NPR / WFDD 2026-01-30	oncology / unsafe treatment advice	include with caution	open source
LHA-077	personal column Lupus News Today 2024-11-25	lupus / appointment preparation	include	open source
LHA-078	public post Reddit r/cfs 2026-04-24	chronic fatigue / dysautonomia / endocrine	include with caution	open source
LHA-079	public comment Hacker News 2023-04-29	unknown chronic illness / possible autoimmune	include	open source
LHA-080	personal essay Medium 2025-06-15	trauma / anxiety / chronic symptoms / MCAS hypothesis	include with caution	open source
LHA-081	public app review Apple App Store 2026-03-ish	chronic illness / invisible illness	candidate	open source
LHA-082	public thread comment Reddit r/ChatGPT 2025-08-25	long COVID / mast cell disorder / anaphylaxis	candidate	open source
LHA-083	public thread comment Reddit r/ChatGPT 2025-08-25	allergies / medication adverse effects	include with caution	open source
LHA-084	public thread comment Reddit r/ChatGPT 2025-08-25	hypertension / endocrine / primary aldosteronism	candidate	open source
LHA-085	public thread comment Reddit r/ChatGPT 2025-08-25	long COVID / chronic fatigue syndrome	include	open source

ID	Source / platform	Health domain	Status	URL
LHA-086	public post Reddit r/OpenAI 2026-04-17	sleep medicine / possible UARS	include	open source
LHA-087	public thread comment Reddit r/OpenAI 2026-04-18	sleep medicine / UARS	candidate	open source
LHA-088	public post Reddit r/ChatGPT 2026-03-14	ENT / Frey's syndrome	include with caution	open source
LHA-089	public thread comment Reddit r/ChatGPT 2026-03-14	respiratory / laryngospasm	candidate	open source
LHA-090	public thread comment Reddit r/ChatGPT 2026-03-14	endocrine / thyroid nodules / imaging report	candidate	open source
LHA-091	public thread comment Reddit r/ChatGPT 2026-03-14	medical report interpretation	include	open source
LHA-092	public thread comment Reddit r/ChatGPT 2025-12-18	obstetrics / HELLP syndrome / liver abnormalities	candidate	open source
LHA-093	public thread comment Reddit r/ChatGPT 2025-12-18	colorectal / oncology	candidate	open source
LHA-094	public thread comment Reddit r/ChatGPT 2025-12-18	neurology / symptom tracking	candidate	open source
LHA-095	public thread comment Reddit r/ChatGPT 2025-12-18	orthopedics / cervical spine	candidate	open source
LHA-096	patient forum Mayo Clinic Connect 2026-01-12	oncology / bladder cancer	candidate	open source
LHA-097	patient forum Mayo Clinic Connect 2025-12-14	oncology / prostate cancer / pathology	include with caution	open source
LHA-098	patient forum Mayo Clinic Connect 2025-ish	osteoporosis / cancer history / diet	include	open source
LHA-099	patient forum Mayo Clinic Connect 2025-ish	hematology / myeloproliferative neoplasm	candidate	open source
LHA-100	patient forum Mayo Clinic Connect 2025-11-20	oncology / biopsy decision support	include	open source